Next Steps

- See if similar techniques can be employed to study matrix-vector product lower bounds for

- Current algorithms use $O(k/\sqrt{\varepsilon})$ matrix-vector products

- Tight for $k = 1$ [Bakshi, Narayanan '23]

- Extending it to all k?

$$\|A - B\|_2 \leq (1 + \varepsilon)\sigma_{k+1}(A)$$

Next Steps

- See if similar techniques can be employed to study matrix-vector product lower bounds for

$$\|A - B\|_2 \leq (1 + \varepsilon)\sigma_{k+1}(A)$$

- Current algorithms use $O(k/\sqrt{\varepsilon})$ matrix-vector products
  - Tight for $k = 1$ [Bakshi, Narayanan '23]
  - Extending it to all k?

Turnstile Streaming

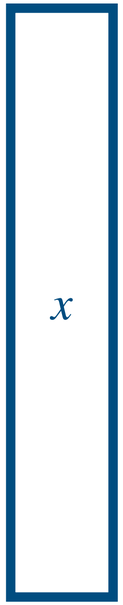

Updates: (i_1, Δ_1) (i_2, Δ_2) $\ldots$ (i_m, Δ_m)